


# High-Performance Multi-Threaded Chat Server

Complete Project Documentation  
& Interview Preparation Guide



## Key Technologies

<b>C++17</b>	<b>IOCP</b>	<b>Thread Pool</b>
Systems Programming	Async I/O	Concurrency

Pranav Kashyap

February 6, 2026

# Contents

<b>1</b>	<b>Project Overview</b>	<b>3</b>
1.1	What is this Project? . . . . .	3
1.2	Key Features . . . . .	3
<b>2</b>	<b>Why C++? The Language Choice</b>	<b>3</b>
2.1	Primary Reasons . . . . .	3
2.1.1	1. Performance is Critical . . . . .	3
2.1.2	2. Direct Hardware & OS Access . . . . .	3
2.1.3	3. Low-Level Socket Control . . . . .	3
2.1.4	4. Thread Control . . . . .	4
2.1.5	5. Memory Efficiency . . . . .	4
2.2	Why NOT Other Languages? . . . . .	4
<b>3</b>	<b>What Makes This Project Different?</b>	<b>4</b>
3.1	Comparison: Normal Chat App vs. This Project . . . . .	4
3.2	Key Differentiators . . . . .	4
3.2.1	1. IOCP – The Secret Weapon . . . . .	4
3.2.2	2. Thread Pool Architecture . . . . .	5
3.2.3	3. Production-Ready Features . . . . .	5
3.2.4	4. Modular Design . . . . .	5
<b>4</b>	<b>Architecture Deep Dive</b>	<b>5</b>
4.1	System Architecture . . . . .	5
4.2	Data Flow . . . . .	6
<b>5</b>	<b>Core Components Explained</b>	<b>6</b>
5.1	Component 1: IOCP Server . . . . .	6
5.2	Component 2: Thread Pool . . . . .	7
5.3	Component 3: Connection Manager . . . . .	7
5.4	Component 4: Chat Room Manager . . . . .	7
5.5	Component 5: Message Store . . . . .	8
<b>6</b>	<b>Technical Concepts Used</b>	<b>8</b>
<b>7</b>	<b>Complete Interview Q&amp;A</b>	<b>8</b>
7.1	Category A: Project Overview Questions . . . . .	8
7.1.1	Q1: What is this project? . . . . .	8
7.1.2	Q2: Why did you choose C++ for this project? . . . . .	8
7.1.3	Q3: What makes this different from a basic chat application? . . . . .	9
7.2	Category B: Architecture & Design Questions . . . . .	9
7.2.1	Q4: Explain the architecture of your system. . . . .	9
7.2.2	Q5: What is IOCP and why did you use it? . . . . .	9
7.2.3	Q6: How does your Thread Pool work? . . . . .	9
7.2.4	Q7: Why a Thread Pool instead of creating threads per request? . . . . .	10
7.2.5	Q8: How do you handle concurrent access to shared data? . . . . .	10
7.3	Category C: Networking Questions . . . . .	10
7.3.1	Q9: Explain the difference between blocking and non-blocking I/O. . . . .	10

7.3.2	Q10: How many clients can your server handle? . . . . .	10
7.3.3	Q11: How do you handle client disconnections? . . . . .	11
7.4	Category D: Threading & Concurrency Questions . . . . .	11
7.4.1	Q12: What is a race condition? How do you prevent them? . . . .	11
7.4.2	Q13: What is a deadlock? How do you prevent it? . . . . .	11
7.4.3	Q14: What is the difference between a mutex and a semaphore? .	11
7.5	Category E: Data Structures Questions . . . . .	12
7.5.1	Q15: What data structures do you use and why? . . . . .	12
7.5.2	Q16: Why <code>unordered_map</code> over <code>map</code> ? . . . . .	12
7.6	Category F: Feature-Specific Questions . . . . .	12
7.6.1	Q17: How does rate limiting work? . . . . .	12
7.6.2	Q18: How is message persistence implemented? . . . . .	12
7.6.3	Q19: How do chat rooms work? . . . . .	13
7.7	Category G: Error Handling Questions . . . . .	13
7.7.1	Q20: How do you handle errors? . . . . .	13
7.7.2	Q21: What happens if the server crashes? . . . . .	13
7.8	Category H: Performance Questions . . . . .	13
7.8.1	Q22: What is the time complexity of your operations? . . . . .	13
7.8.2	Q23: How would you benchmark this server? . . . . .	14
7.9	Category I: Security Questions . . . . .	14
7.9.1	Q24: What security features does your server have? . . . . .	14
7.9.2	Q25: What security improvements would you add? . . . . .	14
7.10	Category J: Code Quality Questions . . . . .	14
7.10.1	Q26: How is your code organized? . . . . .	14
7.10.2	Q27: What design patterns do you use? . . . . .	15
7.11	Category K: Improvement Questions . . . . .	15
7.11.1	Q28: What would you improve in this project? . . . . .	15
7.11.2	Q29: How would you scale this to millions of users? . . . . .	15
7.11.3	Q30: Can this run on Linux? . . . . .	15
<b>8</b>	<b>Code Walkthrough</b>	<b>16</b>
8.1	Main Server Flow . . . . .	16
8.2	Message Handling Flow . . . . .	16
<b>9</b>	<b>Challenges &amp; Solutions</b>	<b>17</b>
<b>10</b>	<b>Future Improvements</b>	<b>17</b>
10.1	Short Term . . . . .	17
10.2	Long Term . . . . .	17
	<b>Summary: Key Talking Points</b>	<b>17</b>

# 1 Project Overview

## 1.1 What is this Project?

A **high-performance, real-time chat server** built in C++ for Windows that can handle **1000+ concurrent connections** using advanced OS-level optimizations.

## 1.2 Key Features

Feature	Description
Multi-threaded Architecture	Thread Pool with worker threads for parallel processing
IOCP (I/O Completion Ports)	Windows' most efficient async I/O mechanism
Multiple Chat Rooms	Users can create, join, and switch between rooms
Private Messaging	Whisper functionality for 1-to-1 communication
Message Persistence	Chat history saved to files, retrievable later
Rate Limiting	Protection against spam and DDoS attacks
Admin Controls	Kick, ban, and mute capabilities

# 2 Why C++? The Language Choice

## 2.1 Primary Reasons

### 2.1.1 1. Performance is Critical

Language	Messages/Second
C++	~1,000,000
Java	~300,000
Python	~50,000

- Chat servers require **microsecond-level latency**
- C++ provides **zero-cost abstractions** – you only pay for what you use
- No garbage collection pauses that could cause message delays

### 2.1.2 2. Direct Hardware & OS Access

- **IOCP (I/O Completion Ports)** is a Windows kernel feature
- C++ can directly call Windows APIs without wrappers
- Control over memory layout, cache optimization

### 2.1.3 3. Low-Level Socket Control

- Direct access to Winsock2 API
- Fine-grained control over TCP/IP parameters
- Buffer management at byte level

### 2.1.4 4. Thread Control

- Precise control over thread creation, affinity, and priority
- No runtime overhead from managed threading
- Direct use of OS synchronization primitives

### 2.1.5 5. Memory Efficiency

- Manual memory management = predictable performance
- No GC pauses during high-traffic periods
- Optimal memory layout for cache performance

## 2.2 Why NOT Other Languages?

Language	Why Not Suitable
Python	Too slow for high-throughput, GIL limits threading
Java	GC pauses, JVM startup overhead, no direct IOCP access
Node.js	Single-threaded event loop, not true parallelism
Go	Good alternative, but less control over OS primitives
Rust	Viable, but more complex syntax, smaller ecosystem

## 3 What Makes This Project Different?

### 3.1 Comparison: Normal Chat App vs. This Project

Aspect	Normal Chat App	This Project
I/O Model	Blocking sockets	IOCP (Async I/O)
Threading	1 thread per client	Thread Pool (fixed)
Scalability	~100 clients	~1000+ clients
Language	Python/Node.js	C++
OS Integration	Generic sockets	Windows kernel IOCP
Complexity	Basic	Production-grade

### 3.2 Key Differentiators

#### 3.2.1 1. IOCP – The Secret Weapon

<p>Traditional: 1 thread waiting per client = 1000 clients = 1000 threads</p> <p>IOCP: 4 worker threads can handle 1000+ clients efficiently</p>
--

### 3.2.2 2. Thread Pool Architecture

- **No thread thrashing:** Fixed number of threads regardless of connections
- **Task queue:** Work items are queued, not thread-spawned
- **CPU-bound optimization:** Thread count = CPU cores

### 3.2.3 3. Production-Ready Features

- Rate limiting (anti-spam, anti-DDoS)
- Connection timeouts
- Message persistence
- Admin moderation tools

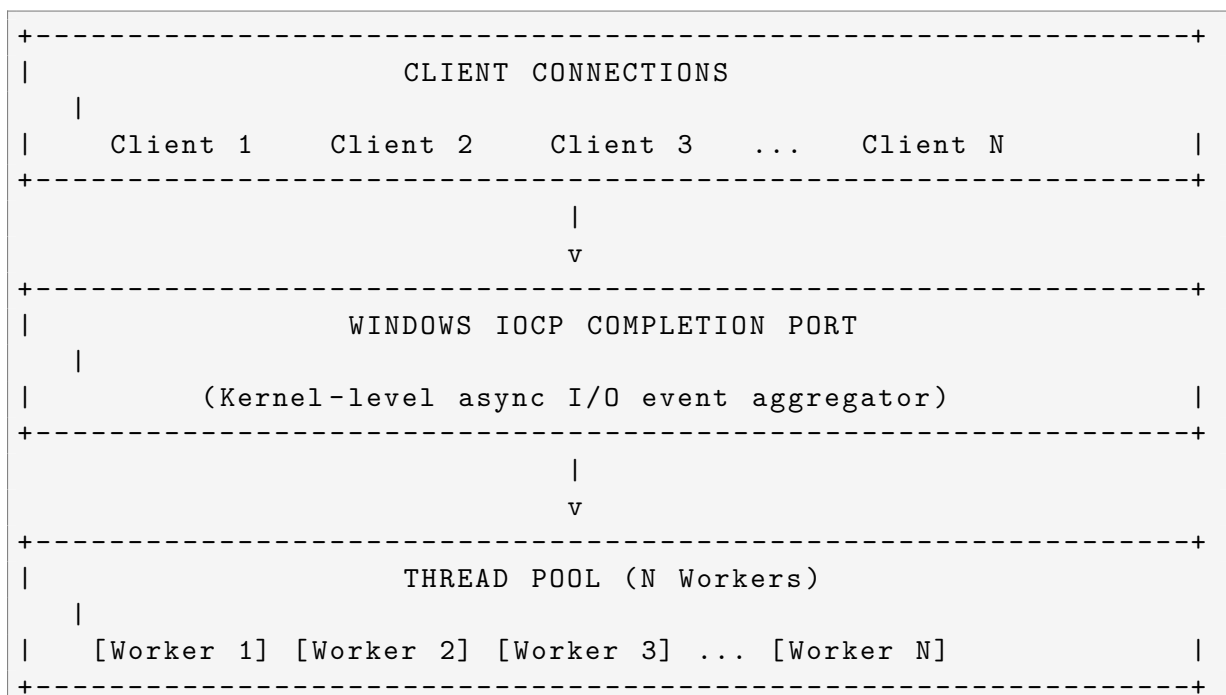
### 3.2.4 4. Modular Design

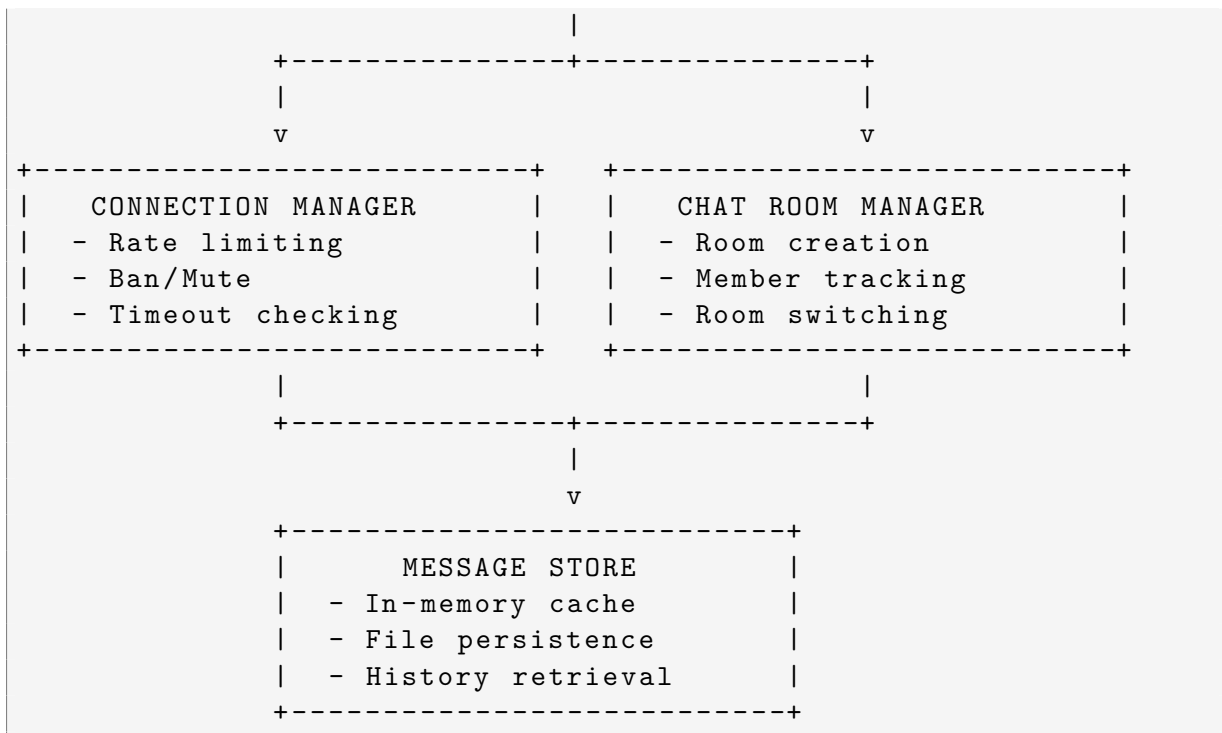
Each component is isolated:

- `ConnectionManager` – Security
- `ChatRoomManager` – Room logic
- `MessageStore` – Persistence
- `IOCPServer` – Networking
- `ThreadPool` – Concurrency

## 4 Architecture Deep Dive

## 4.1 System Architecture





## 4.2 Data Flow

1. Client sends message
2. IOCP receives I/O completion event
3. Worker thread picks up the event
4. Message parsed → Command or Chat?
5. If Command: Execute and respond
6. If Chat: Store message + Broadcast to room members

## 5 Core Components Explained

### 5.1 Component 1: IOCP Server

**File:** iocp\_server.h/cpp

**Purpose:** Handle all network I/O asynchronously

```

1 // Creates completion port
2 HANDLE iocp = CreateIoCompletionPort(...);
3
4 // Worker threads wait for events
5 GetQueuedCompletionStatus(iocp, &bytes, &key, &overlapped, INFINITE);
6
7 // Process the I/O operation
8 HandleRead(socket, buffer, bytes);
  
```

Why IOCP?

- Kernel aggregates all I/O events
- Multiple threads can wait on single port
- No polling overhead
- Scales  $O(1)$  with connection count

## 5.2 Component 2: Thread Pool

**File:** thread\_pool.h/cpp

**Purpose:** Manage worker threads efficiently

```

1 class ThreadPool {
2     std::queue<std::function<void()>> tasks;    // Work queue
3     std::vector<std::thread> workers;           // Fixed thread count
4     std::mutex queue_mutex;
5     std::condition_variable condition;
6 };

```

**Benefits:**

- No thread creation/destruction per request
- Prevents context switching overhead
- Bounded resource usage

## 5.3 Component 3: Connection Manager

**File:** connection\_manager.h/cpp

**Purpose:** Security and stability

Feature	Implementation
Rate Limiting	std::deque with timestamps, sliding window
Ban List	std::unordered_set<string> of IPs
Mute List	std::unordered_map<int, time_point>
Timeout	Last activity timestamp tracking

## 5.4 Component 4: Chat Room Manager

**File:** chat\_room.h/cpp

**Purpose:** Manage logical groupings of users

```

1 struct Room {
2     std::string name;
3     std::unordered_set<int> members;    // O(1) lookup
4     int owner_id;
5     bool is_private;
6 };

```



## 5.5 Component 5: Message Store

**File:** `message_store.h/cpp`

**Purpose:** Persistence and history

**Dual Storage:**

1. **In-Memory Cache:** `std::deque` per room (fast retrieval)
2. **File Storage:** Append-only log files (persistence)

## 6 Technical Concepts Used

1. **I/O Completion Ports (IOCP):** Windows kernel feature for async I/O
2. **Thread Pool Pattern:** Pre-created workers, task queue, producer-consumer
3. **Mutex & Lock Guards:** RAII-style thread synchronization
4. **Smart Pointers:** `std::unique_ptr` for automatic cleanup
5. **Async Event-Driven Programming:** Callbacks for events
6. **Rate Limiting Algorithm:** Sliding window with timestamp deque

## 7 Complete Interview Q&A

### 7.1 Category A: Project Overview Questions

#### 7.1.1 Q1: What is this project?

**Answer:** A high-performance, multi-threaded chat server built in C++ for Windows. It uses I/O Completion Ports (IOCP) for async networking and a Thread Pool for efficient task processing. It supports 1000+ concurrent connections, multiple chat rooms, message persistence, and admin moderation tools.

#### 7.1.2 Q2: Why did you choose C++ for this project?

**Answer:**

1. **Performance:** Chat servers need microsecond-level latency. C++ has zero runtime overhead and no garbage collection pauses.
2. **Direct OS Access:** IOCP is a Windows kernel feature. C++ can call Windows APIs directly without wrappers.
3. **Memory Control:** Manual memory management ensures predictable performance under high load.
4. **Threading Control:** Direct access to OS threading primitives for optimal parallelism.

### 7.1.3 Q3: What makes this different from a basic chat application?

**Answer:**

- **I/O Model:** Uses IOCP (async) instead of blocking sockets
- **Threading:** Thread Pool instead of 1-thread-per-client
- **Scalability:** Handles 1000+ clients vs. ~100 for basic apps
- **Production Features:** Rate limiting, persistence, admin tools
- **Architecture:** Modular, component-based design

## 7.2 Category B: Architecture & Design Questions

### 7.2.1 Q4: Explain the architecture of your system.

**Answer:** The system follows a layered architecture:

1. **Network Layer** (IOCP Server): Handles all socket I/O asynchronously
2. **Processing Layer** (Thread Pool): Workers pull tasks from a queue
3. **Business Logic Layer:** Connection Manager, Chat Room Manager, Message Store
4. **Persistence Layer:** File-based message logging

All layers are loosely coupled through clean interfaces.

### 7.2.2 Q5: What is IOCP and why did you use it?

**Answer:** IOCP (I/O Completion Ports) is Windows' most efficient async I/O mechanism. When an I/O operation completes, the kernel posts a notification to a completion port. Multiple worker threads can wait on this port.

**Benefits:**

- Single kernel object handles all sockets
- Automatic load balancing across threads
- Scales  $O(1)$  with connection count
- No polling overhead

### 7.2.3 Q6: How does your Thread Pool work?

**Answer:**

```
1 class ThreadPool {
2     queue<function<void()>> tasks;    // Task queue
3     vector<thread> workers;          // Fixed threads
4     mutex queue_mutex;
5     condition_variable cv;
6 };
```

Workers wait on a condition variable. When a task is added:

1. Lock queue
2. Push task
3. Notify one worker
4. Worker wakes, pops task, executes

#### 7.2.4 Q7: Why a Thread Pool instead of creating threads per request?

**Answer:**

1. **Thread creation is expensive:** ~1ms per thread on Windows
2. **Context switching overhead:** Many threads = CPU wasted on switching
3. **Bounded resources:** Fixed thread count prevents exhaustion
4. **No thread thrashing:** Even under high load, threads are reused

#### 7.2.5 Q8: How do you handle concurrent access to shared data?

**Answer:** Multiple synchronization mechanisms:

1. **Mutexes:** `w32::Mutex` wraps Windows `CRITICAL_SECTION`
2. **Lock Guards:** RAII-style `w32::LockGuard` for automatic unlock
3. **Atomic variables:** `std::atomic<bool>` for simple flags
4. **Per-component locks:** Each manager has its own mutex

### 7.3 Category C: Networking Questions

#### 7.3.1 Q9: Explain the difference between blocking and non-blocking I/O.

**Answer:**

Blocking I/O	Non-Blocking I/O
<code>recv()</code> waits until data arrives	Returns immediately (-1 if no data)
Thread is stuck waiting	Thread can do other work
Simple to code	Requires event loop/polling
Poor scalability	Excellent scalability

IOCP is **async I/O** – even better: the OS notifies you when I/O completes.

#### 7.3.2 Q10: How many clients can your server handle?

**Answer:** Designed for 1000+ concurrent connections. Limitations:

- **Configured limit:** `max_total_connections = 1000`
- **FD limit:** Windows has no per-process socket limit like Linux
- **Memory:** ~100KB per client = 1GB for 10,000 clients
- **IOCP scales O(1):** Kernel handles the multiplexing

### 7.3.3 Q11: How do you handle client disconnections?

**Answer:**

1. **Detection:** `recv()` returns 0 or IOCP posts disconnect event
2. **Cleanup:** Remove from chat room, delete from client name map, update connection count, notify room members
3. **Graceful handling:** Socket shutdown before close

## 7.4 Category D: Threading & Concurrency Questions

### 7.4.1 Q12: What is a race condition? How do you prevent them?

**Answer:** A race condition occurs when multiple threads access shared data concurrently and at least one is writing.

**Prevention in my project:**

1. Mutex locks around critical sections
2. Lock guards for exception-safe locking
3. Atomic operations for simple counters
4. Copying data before returning (return copy, not reference)

### 7.4.2 Q13: What is a deadlock? How do you prevent it?

**Answer:** Deadlock occurs when threads wait for each other's locks.

**Prevention strategies:**

1. **Lock ordering:** Always acquire locks in same order
2. **Lock timeout:** Use `TryLock` with timeout
3. **Fine-grained locking:** Each component has its own mutex
4. **Avoid holding locks:** Copy data, release lock, then process

### 7.4.3 Q14: What is the difference between a mutex and a semaphore?

**Answer:**

Mutex	Semaphore
Binary (locked/unlocked)	Counter (0 to N)
Ownership (only owner unlocks)	No ownership
For mutual exclusion	For resource counting
Example: Data structure access	Example: Connection pool

I use mutexes because I need exclusive access to data structures.

## 7.5 Category E: Data Structures Questions

### 7.5.1 Q15: What data structures do you use and why?

**Answer:**

Data Structure	Use Case	Why
<code>unordered_map</code>	Client names, rooms	$O(1)$ lookup
<code>unordered_set</code>	Room members, banned IPs	$O(1)$ insert/lookup
<code>deque</code>	Message history, timestamps	$O(1)$ front/back ops
<code>vector</code>	Worker threads, client lists	Cache-friendly
<code>queue</code>	Task queue	FIFO ordering

### 7.5.2 Q16: Why `unordered_map` over `map`?

**Answer:**

- `unordered_map`:  $O(1)$  average lookup (hash table)
- `map`:  $O(\log n)$  lookup (red-black tree)

For chat server with frequent lookups by client ID,  $O(1)$  is essential.

## 7.6 Category F: Feature-Specific Questions

### 7.6.1 Q17: How does rate limiting work?

**Answer:** Sliding window algorithm:

```

1 bool AllowMessage(int client_id) {
2     auto now = steady_clock::now();
3     auto& timestamps = client_messages[client_id];
4
5     // Remove old timestamps (>1 minute)
6     while (!timestamps.empty() &&
7           now - timestamps.front() > minutes(1)) {
8         timestamps.pop_front();
9     }
10
11     // Check limit
12     return timestamps.size() < max_messages_per_minute;
13 }
```

This prevents spam while allowing burst of messages.

### 7.6.2 Q18: How is message persistence implemented?

**Answer:** Dual storage approach:

1. **In-Memory Cache:** deque per room with max size (e.g., 100 messages)
2. **File Persistence:** Append-only log files in `./chat_logs/`

### 7.6.3 Q19: How do chat rooms work?

**Answer:**

- Each client is in exactly one room at a time
- Default room: `#general`
- Data structures: `rooms` (map of Room), `client_rooms` (client  $\rightarrow$  room)

Join operation: Leave current room  $\rightarrow$  Add to new room  $\rightarrow$  Update mapping  $\rightarrow$  Notify both rooms

## 7.7 Category G: Error Handling Questions

### 7.7.1 Q20: How do you handle errors?

**Answer:**

1. **Socket errors:** Check return values, gracefully disconnect
2. **Memory allocation:** `std::make_unique` throws on failure
3. **Rate limit violations:** Send error message to client
4. **Component failures:** Continue running, log error

### 7.7.2 Q21: What happens if the server crashes?

**Answer:**

- **Messages:** Persisted to files, can be recovered
- **Connections:** Clients detect disconnect
- **Cleanup:** RAII ensures resources are released
- **Restart:** Server can restart and load history

## 7.8 Category H: Performance Questions

### 7.8.1 Q22: What is the time complexity of your operations?

**Answer:**

Operation	Complexity
Send message	$O(m)$ where $m$ = room members
Join room	$O(1)$ hash table ops
Find user	$O(1)$ hash table lookup
Get history	$O(n)$ where $n$ = requested messages
Rate limit check	$O(1)$ amortized

### 7.8.2 Q23: How would you benchmark this server?

**Answer:**

1. **Connection test:** Connect 1000+ clients simultaneously
2. **Throughput test:** Measure messages/second
3. **Latency test:** Time from send to receive
4. **Memory test:** Monitor memory under load
5. **CPU test:** Profile worker thread utilization

## 7.9 Category I: Security Questions

### 7.9.1 Q24: What security features does your server have?

**Answer:**

1. **Rate limiting:** Prevents spam/DDoS
2. **IP banning:** Permanent block for abusers
3. **User muting:** Temporary silence
4. **Connection limiting:** Max 1000 concurrent
5. **Timeout:** Disconnect idle clients

### 7.9.2 Q25: What security improvements would you add?

**Answer:**

1. **TLS/SSL:** Encrypt all communications
2. **Authentication:** User accounts with passwords
3. **Input validation:** Sanitize all messages
4. **Admin authentication:** Verify admin commands
5. **Rate limit by IP:** Not just by client ID

## 7.10 Category J: Code Quality Questions

### 7.10.1 Q26: How is your code organized?

**Answer:** Modular component-based design:

<code>server.cpp</code>	- Entry point & event handlers
<code>client.cpp</code>	- Client application
<code>iocp_server.h/cpp</code>	- Network layer
<code>thread_pool.h/cpp</code>	- Concurrency
<code>connection_manager</code>	- Security
<code>chat_room</code>	- Room logic
<code>message_store</code>	- Persistence
<code>socketutil</code>	- Socket utilities

<code>win32_compat</code> - Platform compatibility
--

Each component has single responsibility.

### 7.10.2 Q27: What design patterns do you use?

**Answer:**

1. **Singleton-like:** Global component pointers
2. **Observer/Callback:** Event handlers for connect/message
3. **Producer-Consumer:** Thread pool task queue
4. **RAII:** Lock guards, smart pointers
5. **Command Pattern:** Chat commands like #join, #kick

## 7.11 Category K: Improvement Questions

### 7.11.1 Q28: What would you improve in this project?

**Answer:**

1. **SSL/TLS encryption** for security
2. **Database persistence** (SQLite) for better querying
3. **WebSocket support** for browser clients
4. **User authentication** with passwords
5. **File sharing** capability
6. **Load balancing** for multiple server instances

### 7.11.2 Q29: How would you scale this to millions of users?

**Answer:**

1. **Horizontal scaling:** Multiple server instances
2. **Load balancer:** Distribute connections
3. **Message broker:** Redis pub/sub for cross-server messaging
4. **Database cluster:** Distributed message storage
5. **Microservices:** Separate auth, rooms, messaging
6. **CDN:** For file/media sharing

### 7.11.3 Q30: Can this run on Linux?

**Answer:** Currently Windows-only due to IOCP. For cross-platform:

- Replace IOCP with **epoll** (Linux) or **kqueue** (macOS)
- Use cross-platform library like **libuv** or **Boost.Asio**



- Replace Windows thread primitives with `std::thread`

## 8 Code Walkthrough

### 8.1 Main Server Flow

```
1 int main() {
2     // 1. Initialize Winsock
3     InitializeWinsock();
4
5     // 2. Create components
6     g_thread_pool = make_unique<ThreadPool>(cpu_cores);
7     g_connection_manager = make_unique<ConnectionManager>(config);
8     g_chat_rooms = make_unique<ChatRoomManager>();
9     g_message_store = make_unique<MessageStore>(config);
10    g_server = make_unique<IOCPServer>(port, *g_thread_pool);
11
12    // 3. Register event handlers
13    g_server->OnMessage(HandleMessage);
14    g_server->OnConnect(HandleConnect);
15    g_server->OnDisconnect(HandleDisconnect);
16
17    // 4. Start server
18    g_server->Start();
19
20    // 5. Main loop (check timeouts)
21    while (running) {
22        auto timed_out = g_connection_manager->CheckTimeouts(...);
23        for (int id : timed_out) {
24            g_server->DisconnectClient(id);
25        }
26        Sleep(1000);
27    }
28
29    // 6. Cleanup
30    CleanupWinsock();
31 }
```

### 8.2 Message Handling Flow

```
1 void HandleMessage(int client_id, const char* message, int length) {
2     // 1. Parse and trim message
3     string msg(message, length);
4
5     // 2. Check rate limiting
6     if (!g_connection_manager->AllowMessage(client_id)) {
7         SendToClient(client_id, "Slow down!");
8         return;
9     }
10
11    // 3. Check if muted
12    if (g_connection_manager->IsMuted(client_id)) {
13        SendToClient(client_id, "You are muted.");
14        return;
15    }
```

```

16
17 // 4. Command or chat?
18 if (msg[0] == '#') {
19     ProcessCommand(client_id, msg);
20 } else {
21     BroadcastToRoom(client_id, msg);
22 }
23 }

```

## 9 Challenges & Solutions

Challenge	Solution
Thread safety for shared data	Mutex per component + lock guards
Handling 1000+ connections	IOCP + Thread Pool
Preventing spam	Rate limiting with sliding window
Message loss on crash	File persistence
Graceful shutdown	Signal handler + resource cleanup
Cross-version Windows support	win32_compat.h wrapper

## 10 Future Improvements

### 10.1 Short Term

Add SSL/TLS encryption

Implement user authentication

Add file sharing

### 10.2 Long Term

WebSocket support for browsers

Database storage (SQLite/PostgreSQL)

Horizontal scaling with Redis

Mobile client apps

## Summary: Key Talking Points

When presenting this project, emphasize:

1. **Technical Depth:** IOCP, Thread Pools, Async I/O
2. **Scalability:** Designed for 1000+ concurrent connections
3. **Production-Ready Features:** Rate limiting, persistence, admin tools
4. **Code Quality:** Modular design, RAII, proper synchronization

5. **Real-World Applications:** Applicable to gaming, collaboration, social apps

---

*Document created for interview preparation. Contains comprehensive explanations of all technical decisions, architecture choices, and potential interview questions with detailed answers.*